

Image dataset: <https://www.kaggle.com/datasets/hlrhegemony/pokemon-image-dataset?resource=download>

Metainfo csv: <https://github.com/lgreski/pokemonData/blob/master/Pokemon.csv>

You can speed up model inference on colab by changing Hardware accelerator in runtime type to any GPU option

```
!pip install datasets transformers numpy pandas Pillow matplotlib
!pip install torch tqdm scipy
!pip install git+https://github.com/openai/CLIP.git
!pip install plotly umap-learn

Requirement already satisfied: datasets in
/usr/local/lib/python3.12/dist-packages (4.0.0)
Requirement already satisfied: transformers in
/usr/local/lib/python3.12/dist-packages (5.0.0)
Requirement already satisfied: numpy in
/usr/local/lib/python3.12/dist-packages (2.0.2)
Requirement already satisfied: pandas in
/usr/local/lib/python3.12/dist-packages (2.2.2)
Requirement already satisfied: Pillow in
/usr/local/lib/python3.12/dist-packages (11.3.0)
Requirement already satisfied: matplotlib in
/usr/local/lib/python3.12/dist-packages (3.10.0)
Requirement already satisfied: filelock in
/usr/local/lib/python3.12/dist-packages (from datasets) (3.24.2)
Requirement already satisfied: pyarrow>=15.0.0 in
/usr/local/lib/python3.12/dist-packages (from datasets) (18.1.0)
Requirement already satisfied: dill<0.3.9,>=0.3.0 in
/usr/local/lib/python3.12/dist-packages (from datasets) (0.3.8)
Requirement already satisfied: requests>=2.32.2 in
/usr/local/lib/python3.12/dist-packages (from datasets) (2.32.4)
Requirement already satisfied: tqdm>=4.66.3 in
/usr/local/lib/python3.12/dist-packages (from datasets) (4.67.3)
Requirement already satisfied: xxhash in
/usr/local/lib/python3.12/dist-packages (from datasets) (3.6.0)
Requirement already satisfied: multiprocessing<0.70.17 in
/usr/local/lib/python3.12/dist-packages (from datasets) (0.70.16)
Requirement already satisfied: fsspec<=2025.3.0,>=2023.1.0 in
/usr/local/lib/python3.12/dist-packages (from
fsspec[http]<=2025.3.0,>=2023.1.0->datasets) (2025.3.0)
Requirement already satisfied: huggingface-hub>=0.24.0 in
/usr/local/lib/python3.12/dist-packages (from datasets) (1.4.1)
Requirement already satisfied: packaging in
/usr/local/lib/python3.12/dist-packages (from datasets) (26.0)
Requirement already satisfied: pyyaml>=5.1 in
/usr/local/lib/python3.12/dist-packages (from datasets) (6.0.3)
Requirement already satisfied: regex!=2019.12.17 in
/usr/local/lib/python3.12/dist-packages (from transformers)
```

(2025.11.3)

```
Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in
/usr/local/lib/python3.12/dist-packages (from transformers) (0.22.2)
Requirement already satisfied: typer-slim in
/usr/local/lib/python3.12/dist-packages (from transformers) (0.24.0)
Requirement already satisfied: safetensors>=0.4.3 in
/usr/local/lib/python3.12/dist-packages (from transformers) (0.7.0)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.12/dist-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in
/usr/local/lib/python3.12/dist-packages (from pandas) (2025.3)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.12/dist-packages (from matplotlib) (1.3.3)
Requirement already satisfied: cycler>=0.10 in
/usr/local/lib/python3.12/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.12/dist-packages (from matplotlib) (4.61.1)
Requirement already satisfied: kiwisolver>=1.3.1 in
/usr/local/lib/python3.12/dist-packages (from matplotlib) (1.4.9)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.12/dist-packages (from matplotlib) (3.3.2)
Requirement already satisfied: aiohttp!=4.0.0a0,!>4.0.0a1 in
/usr/local/lib/python3.12/dist-packages (from
fsspec[http]<=2025.3.0,>=2023.1.0->datasets) (3.13.3)
Requirement already satisfied: hf-xet<2.0.0,>=1.2.0 in
/usr/local/lib/python3.12/dist-packages (from huggingface-hub>=0.24.0-
>datasets) (1.2.0)
Requirement already satisfied: httpx<1,>=0.23.0 in
/usr/local/lib/python3.12/dist-packages (from huggingface-hub>=0.24.0-
>datasets) (0.28.1)
Requirement already satisfied: shellingham in
/usr/local/lib/python3.12/dist-packages (from huggingface-hub>=0.24.0-
>datasets) (1.5.4)
Requirement already satisfied: typing-extensions>=4.1.0 in
/usr/local/lib/python3.12/dist-packages (from huggingface-hub>=0.24.0-
>datasets) (4.15.0)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2-
>pandas) (1.17.0)
Requirement already satisfied: charset_normalizer<4,>=2 in
/usr/local/lib/python3.12/dist-packages (from requests>=2.32.2-
>datasets) (3.4.4)
Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.12/dist-packages (from requests>=2.32.2-
>datasets) (3.11)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.12/dist-packages (from requests>=2.32.2-
```

```
>datasets) (2.5.0)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.12/dist-packages (from requests>=2.32.2-
>datasets) (2026.1.4)
Requirement already satisfied: typer>=0.24.0 in
/usr/local/lib/python3.12/dist-packages (from typer-slim-
>transformers) (0.24.0)
Requirement already satisfied: aiohappyeyeballs>=2.5.0 in
/usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0,!
=4.0.0a1->fsspec[http]<=2025.3.0,>=2023.1.0->datasets) (2.6.1)
Requirement already satisfied: aiosignal>=1.4.0 in
/usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0,!
=4.0.0a1->fsspec[http]<=2025.3.0,>=2023.1.0->datasets) (1.4.0)
Requirement already satisfied: attrs>=17.3.0 in
/usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0,!
=4.0.0a1->fsspec[http]<=2025.3.0,>=2023.1.0->datasets) (25.4.0)
Requirement already satisfied: frozenlist>=1.1.1 in
/usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0,!
=4.0.0a1->fsspec[http]<=2025.3.0,>=2023.1.0->datasets) (1.8.0)
Requirement already satisfied: multidict<7.0,>=4.5 in
/usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0,!
=4.0.0a1->fsspec[http]<=2025.3.0,>=2023.1.0->datasets) (6.7.1)
Requirement already satisfied: propcache>=0.2.0 in
/usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0,!
=4.0.0a1->fsspec[http]<=2025.3.0,>=2023.1.0->datasets) (0.4.1)
Requirement already satisfied: yarl<2.0,>=1.17.0 in
/usr/local/lib/python3.12/dist-packages (from aiohttp!=4.0.0a0,!
=4.0.0a1->fsspec[http]<=2025.3.0,>=2023.1.0->datasets) (1.22.0)
Requirement already satisfied: anyio in
/usr/local/lib/python3.12/dist-packages (from httpx<1,>=0.23.0-
>huggingface-hub>=0.24.0->datasets) (4.12.1)
Requirement already satisfied: httpcore==1.* in
/usr/local/lib/python3.12/dist-packages (from httpx<1,>=0.23.0-
>huggingface-hub>=0.24.0->datasets) (1.0.9)
Requirement already satisfied: h11>=0.16 in
/usr/local/lib/python3.12/dist-packages (from httpcore==1.*-
>httpx<1,>=0.23.0->huggingface-hub>=0.24.0->datasets) (0.16.0)
Requirement already satisfied: click>=8.2.1 in
/usr/local/lib/python3.12/dist-packages (from typer>=0.24.0->typer-
slim->transformers) (8.3.1)
Requirement already satisfied: rich>=12.3.0 in
/usr/local/lib/python3.12/dist-packages (from typer>=0.24.0->typer-
slim->transformers) (13.9.4)
Requirement already satisfied: annotated-doc>=0.0.2 in
/usr/local/lib/python3.12/dist-packages (from typer>=0.24.0->typer-
slim->transformers) (0.0.4)
Requirement already satisfied: markdown-it-py>=2.2.0 in
/usr/local/lib/python3.12/dist-packages (from rich>=12.3.0-
>typer>=0.24.0->typer-slim->transformers) (4.0.0)
```

```
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in
/usr/local/lib/python3.12/dist-packages (from rich>=12.3.0-
>typer>=0.24.0->typer-slim->transformers) (2.19.2)
Requirement already satisfied: mdurl~=0.1 in
/usr/local/lib/python3.12/dist-packages (from markdown-it-py>=2.2.0-
>rich>=12.3.0->typer>=0.24.0->typer-slim->transformers) (0.1.2)
Requirement already satisfied: torch in
/usr/local/lib/python3.12/dist-packages (2.10.0+cu128)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-
packages (4.67.3)
Requirement already satisfied: scipy in
/usr/local/lib/python3.12/dist-packages (1.16.3)
Requirement already satisfied: filelock in
/usr/local/lib/python3.12/dist-packages (from torch) (3.24.2)
Requirement already satisfied: typing-extensions>=4.10.0 in
/usr/local/lib/python3.12/dist-packages (from torch) (4.15.0)
Requirement already satisfied: setuptools in
/usr/local/lib/python3.12/dist-packages (from torch) (75.2.0)
Requirement already satisfied: sympy>=1.13.3 in
/usr/local/lib/python3.12/dist-packages (from torch) (1.14.0)
Requirement already satisfied: networkx>=2.5.1 in
/usr/local/lib/python3.12/dist-packages (from torch) (3.6.1)
Requirement already satisfied: jinja2 in
/usr/local/lib/python3.12/dist-packages (from torch) (3.1.6)
Requirement already satisfied: fsspec>=0.8.5 in
/usr/local/lib/python3.12/dist-packages (from torch) (2025.3.0)
Requirement already satisfied: cuda-bindings==12.9.4 in
/usr/local/lib/python3.12/dist-packages (from torch) (12.9.4)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.8.93 in
/usr/local/lib/python3.12/dist-packages (from torch) (12.8.93)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.8.90 in
/usr/local/lib/python3.12/dist-packages (from torch) (12.8.90)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.8.90 in
/usr/local/lib/python3.12/dist-packages (from torch) (12.8.90)
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in
/usr/local/lib/python3.12/dist-packages (from torch) (9.10.2.21)
Requirement already satisfied: nvidia-cublas-cu12==12.8.4.1 in
/usr/local/lib/python3.12/dist-packages (from torch) (12.8.4.1)
Requirement already satisfied: nvidia-cufft-cu12==11.3.3.83 in
/usr/local/lib/python3.12/dist-packages (from torch) (11.3.3.83)
Requirement already satisfied: nvidia-curand-cu12==10.3.9.90 in
/usr/local/lib/python3.12/dist-packages (from torch) (10.3.9.90)
Requirement already satisfied: nvidia-cusolver-cu12==11.7.3.90 in
/usr/local/lib/python3.12/dist-packages (from torch) (11.7.3.90)
Requirement already satisfied: nvidia-cusparse-cu12==12.5.8.93 in
/usr/local/lib/python3.12/dist-packages (from torch) (12.5.8.93)
Requirement already satisfied: nvidia-cusparseelt-cu12==0.7.1 in
/usr/local/lib/python3.12/dist-packages (from torch) (0.7.1)
Requirement already satisfied: nvidia-nccl-cu12==2.27.5 in
```

```
/usr/local/lib/python3.12/dist-packages (from torch) (2.27.5)
Requirement already satisfied: nvidia-nvshmem-cu12==3.4.5 in
/usr/local/lib/python3.12/dist-packages (from torch) (3.4.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.8.90 in
/usr/local/lib/python3.12/dist-packages (from torch) (12.8.90)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.8.93 in
/usr/local/lib/python3.12/dist-packages (from torch) (12.8.93)
Requirement already satisfied: nvidia-cufile-cu12==1.13.1.3 in
/usr/local/lib/python3.12/dist-packages (from torch) (1.13.1.3)
Requirement already satisfied: triton==3.6.0 in
/usr/local/lib/python3.12/dist-packages (from torch) (3.6.0)
Requirement already satisfied: cuda-pathfinder~=1.1 in
/usr/local/lib/python3.12/dist-packages (from cuda-bindings==12.9.4->torch) (1.3.4)
Requirement already satisfied: numpy<2.6,>=1.25.2 in
/usr/local/lib/python3.12/dist-packages (from scipy) (2.0.2)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in
/usr/local/lib/python3.12/dist-packages (from sympy>=1.13.3->torch)
(1.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.12/dist-packages (from jinja2->torch) (3.0.3)
Collecting git+https://github.com/openai/CLIP.git
  Cloning https://github.com/openai/CLIP.git to /tmp/pip-req-build-k0kc4keg
    Running command git clone --filter=blob:none --quiet
https://github.com/openai/CLIP.git /tmp/pip-req-build-k0kc4keg
  Resolved https://github.com/openai/CLIP.git to commit
ded190a052fdf4585bd685cee5bc96e0310d2c93
  Preparing metadata (setup.py) ... 
Requirement already satisfied: ftfy in
/usr/local/lib/python3.12/dist-packages (from clip==1.0) (6.3.1)
Requirement already satisfied: packaging in
/usr/local/lib/python3.12/dist-packages (from clip==1.0) (26.0)
Requirement already satisfied: regex in
/usr/local/lib/python3.12/dist-packages (from clip==1.0) (2025.11.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-
packages (from clip==1.0) (4.67.3)
Requirement already satisfied: torch in
/usr/local/lib/python3.12/dist-packages (from clip==1.0)
(2.10.0+cu128)
Requirement already satisfied: torchvision in
/usr/local/lib/python3.12/dist-packages (from clip==1.0)
(0.25.0+cu128)
Requirement already satisfied: wcwidth in
/usr/local/lib/python3.12/dist-packages (from ftfy->clip==1.0) (0.6.0)
Requirement already satisfied: filelock in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(3.24.2)
Requirement already satisfied: typing-extensions>=4.10.0 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
```

```
(4.15.0)
Requirement already satisfied: setuptools in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(75.2.0)
Requirement already satisfied: sympy>=1.13.3 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(1.14.0)
Requirement already satisfied: networkx>=2.5.1 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(3.6.1)
Requirement already satisfied: jinja2 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(3.1.6)
Requirement already satisfied: fsspec>=0.8.5 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(2025.3.0)
Requirement already satisfied: cuda-bindings==12.9.4 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(12.9.4)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.8.93 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(12.8.93)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.8.90 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(12.8.90)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.8.90 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(12.8.90)
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(9.10.2.21)
Requirement already satisfied: nvidia-cublas-cu12==12.8.4.1 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(12.8.4.1)
Requirement already satisfied: nvidia-cufft-cu12==11.3.3.83 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(11.3.3.83)
Requirement already satisfied: nvidia-curand-cu12==10.3.9.90 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(10.3.9.90)
Requirement already satisfied: nvidia-cusolver-cu12==11.7.3.90 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(11.7.3.90)
Requirement already satisfied: nvidia-cusparse-cu12==12.5.8.93 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(12.5.8.93)
Requirement already satisfied: nvidia-cusparseelt-cu12==0.7.1 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(0.7.1)
```

```
Requirement already satisfied: nvidia-nccl-cu12==2.27.5 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(2.27.5)
Requirement already satisfied: nvidia-nvshmem-cu12==3.4.5 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(3.4.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.8.90 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(12.8.90)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.8.93 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(12.8.93)
Requirement already satisfied: nvidia-cufile-cu12==1.13.1.3 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(1.13.1.3)
Requirement already satisfied: triton==3.6.0 in
/usr/local/lib/python3.12/dist-packages (from torch->clip==1.0)
(3.6.0)
Requirement already satisfied: cuda-pathfinder~=1.1 in
/usr/local/lib/python3.12/dist-packages (from cuda-bindings==12.9.4-
>torch->clip==1.0) (1.3.4)
Requirement already satisfied: numpy in
/usr/local/lib/python3.12/dist-packages (from torchvision->clip==1.0)
(2.0.2)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in
/usr/local/lib/python3.12/dist-packages (from torchvision->clip==1.0)
(11.3.0)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in
/usr/local/lib/python3.12/dist-packages (from sympy>=1.13.3->torch-
>clip==1.0) (1.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.12/dist-packages (from jinja2->torch-
>clip==1.0) (3.0.3)
Requirement already satisfied: plotly in
/usr/local/lib/python3.12/dist-packages (5.24.1)
Requirement already satisfied: umap-learn in
/usr/local/lib/python3.12/dist-packages (0.5.11)
Requirement already satisfied: tenacity>=6.2.0 in
/usr/local/lib/python3.12/dist-packages (from plotly) (9.1.4)
Requirement already satisfied: packaging in
/usr/local/lib/python3.12/dist-packages (from plotly) (26.0)
Requirement already satisfied: numpy>=1.23 in
/usr/local/lib/python3.12/dist-packages (from umap-learn) (2.0.2)
Requirement already satisfied: scipy>=1.3.1 in
/usr/local/lib/python3.12/dist-packages (from umap-learn) (1.16.3)
Requirement already satisfied: scikit-learn>=1.6 in
/usr/local/lib/python3.12/dist-packages (from umap-learn) (1.6.1)
Requirement already satisfied: numba>=0.51.2 in
/usr/local/lib/python3.12/dist-packages (from umap-learn) (0.60.0)
```

```

Requirement already satisfied: pynndescent>=0.5 in
/usr/local/lib/python3.12/dist-packages (from umap-learn) (0.6.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-
packages (from umap-learn) (4.67.3)
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in
/usr/local/lib/python3.12/dist-packages (from numba>=0.51.2->umap-
learn) (0.43.0)
Requirement already satisfied: joblib>=0.11 in
/usr/local/lib/python3.12/dist-packages (from pynndescent>=0.5->umap-
learn) (1.5.3)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/usr/local/lib/python3.12/dist-packages (from scikit-learn>=1.6->umap-
learn) (3.6.0)

from datasets import load_dataset
from transformers import CLIPProcessor, CLIPModel
import numpy as np
import pandas as pd
from glob import glob
from PIL import Image
import matplotlib.pyplot as plt
import clip
import torch
from tqdm import tqdm
from scipy.special import softmax
import plotly.express as px
import plotly.graph_objects as go
from sklearn.manifold import TSNE

print(f"GPU device: {torch.cuda.get_device_name(0)}")

GPU device: NVIDIA A100-SXM4-40GB

# load csv file and image paths to construct pokedex, use
type_to_load=None to load all types, else use a list of types 1 to
load
def construct_pokedex(csv_path='Pokemon.csv', image_dir='./images/',
type_to_load=None):
    pokedex = pd.read_csv(csv_path)
    image_paths = []

    for pokemon_name in pokedex["Name"]:
        imgs = glob(f"{image_dir}/{pokemon_name}/0.jpg")
        if len(imgs) > 0:
            image_paths.append(imgs[0])
        else:
            image_paths.append(None)

    pokedex["image_path"] = image_paths

```

```

pokedex =
pokedex[pokedex["image_path"].notna()].reset_index(drop=True)

# only keep pokemon with distinct id
ids, id_counts = np.unique(pokedex["ID"], return_counts=True)
ids, id_counts = np.array(ids), np.array(id_counts)
keep_ids = ids[id_counts == 1]

pokedex =
pokedex[pokedex["ID"].isin(keep_ids)].reset_index(drop=True)
pokedex["Type2"] = pokedex["Type2"].str.strip()
if type_to_load is not None:
    pokedex =
pokedex[pokedex["Type1"].isin(type_to_load)].reset_index(drop=True)
return pokedex

def load_clip_model():
    device = "cuda" if torch.cuda.is_available() else "cpu"
    model, preprocess = clip.load("ViT-L/14", device=device)
    return model, preprocess, device

# Updated this function to have batch process 32 at once to get speedup
def clip_inference_image(model, preprocess, image_paths, device,
batch_size=32):

    image_embeddings = []

    if device == "cuda":
        effective_batch_size = batch_size
    else:
        effective_batch_size = min(batch_size, 8)

    with torch.no_grad():
        for i in tqdm(range(0, len(image_paths),
effective_batch_size), desc="Processing images"):
            batch_paths = image_paths[i:i+effective_batch_size]

            batch_images = []
            for img_path in batch_paths:
                img = Image.open(img_path)
                img_preprocessed = preprocess(img)
                batch_images.append(img_preprocessed)

            batch_tensor = torch.stack(batch_images).to(device)

            batch_embeddings =
model.encode_image(batch_tensor).detach().cpu().numpy()
            image_embeddings.append(batch_embeddings)

```

```

        image_embeddings = np.concatenate(image_embeddings, axis=0)
        image_embeddings /= np.linalg.norm(image_embeddings, axis=-1,
keepdims=True)
        return image_embeddings

def clip_inference_text(model, preprocess, texts, device):
    with torch.no_grad():
        text_embeddings =
model.encode_text(clip.tokenize(texts).to(device)).detach().cpu().numpy()
        text_embeddings /= np.linalg.norm(text_embeddings, axis=-1,
keepdims=True)
        return text_embeddings

def compute_similarity_text_to_image(image_embeddings,
text_embeddings):
    similarity = softmax((100.0 * image_embeddings @
text_embeddings.T), axis=-1)
    return similarity

def compute_similarity_image_to_text(image_embeddings,
text_embeddings):
    similarity = softmax((100.0 * image_embeddings @
text_embeddings.T), axis=0)
    return similarity

def umap_projection(image_embeddings, n_neighbors=15, min_dist=0.1,
metric='cosine'):
    distance_matrix = np.zeros((image_embeddings.shape[0],
image_embeddings.shape[0]))
    for i in range(image_embeddings.shape[0]):
        for j in range(image_embeddings.shape[0]):
            if i == j:
                distance_matrix[i, j] = 1
            else:
                distance_matrix[i, j] = np.dot(image_embeddings[i],
image_embeddings[j])
    distance_matrix = 1 - distance_matrix
    reducer = TSNE(n_components=2, metric="precomputed",
init="random", random_state=42)
    visualization_data = reducer.fit_transform(distance_matrix)
    return visualization_data

# =====
# Qwen3-VL-2B-Instruct helper
# =====

# If you get import/version issues, install transformers from source:
# pip install -U git+https://github.com/huggingface/transformers

```

```

import os
import torch
from PIL import Image
from transformers import Qwen3VLForConditionalGeneration,
AutoProcessor

def load_qwen3_vl(model_id="Qwen/Qwen3-VL-2B-Instruct"):
    """
    Load Qwen3-VL model + processor.
    Works on GPU if available (recommended).
    """
    model = Qwen3VLForConditionalGeneration.from_pretrained(
        model_id,
        torch_dtype="auto",
        device_map="auto",
    ).eval()

    processor = AutoProcessor.from_pretrained(model_id)
    return model, processor

@torch.no_grad()
def qwen_vl_infer_one(model, processor, image_path, prompt,
max_new_tokens=128):
    """
    Run Qwen3-VL on a single image + prompt.
    Returns the generated text.
    """
    assert os.path.exists(image_path), f"Image not found: {image_path}"

    messages = [
        {
            "role": "user",
            "content": [
                {"type": "image", "image": image_path},
                {"type": "text", "text": prompt},
            ],
        }
    ]

    # Build model inputs using Qwen chat template
    inputs = processor.apply_chat_template(
        messages,
        tokenize=True,
        add_generation_prompt=True,
        return_dict=True,
        return_tensors="pt",
    ).to(model.device)

```

```

# Deterministic decoding
output_ids = model.generate(
    **inputs,
    max_new_tokens=max_new_tokens,
    do_sample=False,
)

# Remove the prompt tokens from the output before decoding
output_ids_trimmed = [out[len(inp):] for inp, out in
zip(inputs.input_ids, output_ids)]

text = processor.batch_decode(
    output_ids_trimmed,
    skip_special_tokens=True,
    clean_up_tokenization_spaces=False,
)[0]

return text.strip()

# Question 20: Text queries for Pokemon types

import os
image_dir = '/content/drive/MyDrive/images'
csv_path = '/content/Pokemon.csv'

pokedex = construct_pokedex(csv_path=csv_path, image_dir=image_dir,
type_to_load=None)

if len(pokedex) == 0:
    raise ValueError("Count not find Pokemon data.")

print(f"Loaded {len(pokedex)} Pokemon")

# CLIP model
clip_model, clip_preprocess, device = load_clip_model()
print(f"CLIP model loaded on {device}")

# Get image embeddings all of the Pokemon
image_paths = pokedex['image_path'].tolist()
print("Compute image embeddings")
image_embeddings = clip_inference_image(clip_model, clip_preprocess,
image_paths, device)
print(image_embeddings.shape)

Loaded 751 Pokemon
CLIP model loaded on cuda
Computing image embeddings...

Processing images: 100%|██████████| 24/24 [00:06<00:00, 3.74it/s]

```

```
Image embeddings shape: (751, 768)
```

```
type_queries = {
    'Bug': ['type: Bug', 'Bug type Pokemon', 'Pokemon with Bug abilities', 'Bug Pokemon'],
    'Fire': ['type: Fire', 'Fire type Pokemon', 'Pokemon with Fire abilities', 'Fire Pokemon'],
    'Grass': ['type: Grass', 'Grass type Pokemon', 'Pokemon with Grass abilities', 'Grass Pokemon'],
    'Dark': ['type: Dark', 'Dark type Pokemon', 'Pokemon with Dark abilities', 'Dark Pokemon', 'Pokemon Dark', 'Pokemon of Dark'],
    'Dragon': ['type: Dragon', 'Dragon type Pokemon', 'Pokemon with Dragon abilities', 'Dragon Pokemon', 'Pokemon Dragon', 'Pokemon of Dragon']
}

def test_query_template(query_templates, type_name, image_embeddings, pokedex, top_k=5):
    """Test different query templates and return the best one with top-k results"""
    best_query = None
    best_scores = None
    best_indices = None
    best_accuracy = 0

    for query_template in query_templates:
        text_emb = clip_inference_text(clip_model, clip_preprocess,
[query_template], device)

        similarities = image_embeddings @ text_emb.T
        similarities = similarities.flatten()

        top_indices = np.argsort(similarities)[::-1][:top_k]
        correct_count = sum(1 for idx in top_indices if
pokedex.iloc[idx]['Type1'] == type_name)
        accuracy = correct_count / top_k

        if accuracy > best_accuracy:
            best_accuracy = accuracy
            best_query = query_template
            best_scores = similarities[top_indices]
            best_indices = top_indices

    return best_query, best_scores, best_indices, best_accuracy

results_q20 = {}
for type_name in ['Bug', 'Fire', 'Grass', 'Dark', 'Dragon']:
    best_query, best_scores, best_indices, accuracy =
```

```

test_query_template(
    type_queries[type_name], type_name, image_embeddings, pokedex,
top_k=5
)
results_q20[type_name] = {
    'query': best_query,
    'indices': best_indices,
    'scores': best_scores,
    'pokemon': [pokedex.iloc[idx] for idx in best_indices]
}
print(f"{type_name}: Best query = '{best_query}', Accuracy = {accuracy:.2%}")

Bug: Best query = 'Bug type Pokemon', Accuracy = 100.00%
Fire: Best query = 'type: Fire', Accuracy = 100.00%
Grass: Best query = 'Pokemon with Grass abilities', Accuracy = 100.00%
Dark: Best query = 'type: Dark', Accuracy = 60.00%
Dragon: Best query = 'Pokemon of Dragon', Accuracy = 60.00%

# Pics for top 5 Pokemon in Bug, Fire, Grass
fig = plt.figure(figsize=(15, 9))

for idx, type_name in enumerate(['Bug', 'Fire', 'Grass']):
    result = results_q20[type_name]

    for i, pokemon_idx in enumerate(result['indices']):
        pokemon = pokedex.iloc[pokemon_idx]
        img_path = pokemon['image_path']

        img = Image.open(img_path)
        ax_sub = plt.subplot(3, 5, idx * 5 + i + 1)
        ax_sub.imshow(img)
        ax_sub.axis('off')
        title = f"{pokemon['Name']}\n{pokemon['Type1']}"
        if pd.notna(pokemon['Type2']) and pokemon['Type2'] != '':
            title += f"/{pokemon['Type2']}"
        ax_sub.set_title(title, fontsize=8)

    fig.text(0.5, 0.98 - idx * 0.33, f"Query: {result['query']}",
             ha='center', fontsize=12, weight='bold')

plt.tight_layout()
plt.savefig('question20_1.png', dpi=150, bbox_inches='tight')
plt.show()

print("\nEffectiveness Assessment:")
for type_name in ['Bug', 'Fire', 'Grass']:
    result = results_q20[type_name]
    correct_count = sum(1 for pokemon in result['pokemon'] if
pokemon['Type1'] == type_name)

```

```
print(f"{type_name}: {correct_count}/5 top results have correct Type1")
```



Effectiveness Assessment:

Bug: 5/5 top results have correct Type1
 Fire: 5/5 top results have correct Type1
 Grass: 5/5 top results have correct Type1

```
# Pics for top 5 Pokémons in Dark and Dragon
fig = plt.figure(figsize=(15, 6))

for idx, type_name in enumerate(['Dark', 'Dragon']):
    result = results_q20[type_name]

    for i, pokemon_idx in enumerate(result['indices']):
        pokemon = pokedex.iloc[pokemon_idx]
        img_path = pokemon['image_path']

        img = Image.open(img_path)
        ax_sub = plt.subplot(2, 5, idx * 5 + i + 1)
        ax_sub.imshow(img)
        ax_sub.axis('off')
        title = f'{pokemon["Name"]}\n{pokemon["Type1"]}'
        if pd.notna(pokemon['Type2']) and pokemon['Type2'] != '':
            title += f' / {pokemon["Type2"]}'
```

```

        ax_sub.set_title(title, fontsize=8)

    fig.text(0.5, 0.98 - idx * 0.5, f"Query: {result['query']}",
             ha='center', fontsize=12, weight='bold')

plt.tight_layout()
plt.savefig('question20_2.png', dpi=150, bbox_inches='tight')
plt.show()

# Could not get above 3/5 with any keywords for Dark and Dragon :(

```

Effectiveness Assessment:

```

print("\nEffectiveness Assessment:")
for type_name in ['Dark', 'Dragon']:
    result = results_q20[type_name]
    correct_count = sum(1 for pokemon in result['pokemon'] if
pokemon['Type1'] == type_name)
    print(f"{type_name}: {correct_count}/5 top results have correct
Type1")

```



Effectiveness Assessment:
Dark: 3/5 top results have correct Type1
Dragon: 3/5 top results have correct Type1

Question 21: Randomly select 10 Pokémons and find their most relevant types

```

all_types = set()
for _, row in pokedex.iterrows():
    if pd.notna(row['Type1']):
        all_types.add(row['Type1'])
    if pd.notna(row['Type2']) and row['Type2'] != '':
        all_types.add(row['Type2'])

```

```

all_types = sorted(list(all_types))
print(f"Found {len(all_types)} unique types: {all_types}")

type_queries_q21 = [f"type: {t}" for t in all_types]

print("Computing text embeddings for all types")
text_embeddings_all_types = clip_inference_text(clip_model,
clip_preprocess, type_queries_q21, device)
print(f"Text embeddings shape: {text_embeddings_all_types.shape}")

np.random.seed(42)
random_indices = np.random.choice(len(pokedex), size=10,
replace=False)
random_pokemon = pokedex.iloc[random_indices]

for idx, pokemon_idx in enumerate(random_indices):
    pokemon = pokedex.iloc[pokemon_idx]
    print(f"{idx+1}. {pokemon['Name']} - Type1: {pokemon['Type1']},
Type2: {pokemon['Type2']}")

Found 18 unique types: ['Bug', 'Dark', 'Dragon', 'Electric', 'Fairy',
'Fighting', 'Fire', 'Flying', 'Ghost', 'Grass', 'Ground', 'Ice',
'Normal', 'Poison', 'Psychic', 'Rock', 'Steel', 'Water']
Computing text embeddings for all types
Text embeddings shape: (18, 768)
1. Sawsbuck - Type1: Normal, Type2: Grass
2. Croagunk - Type1: Poison, Type2: Fighting
3. Hoppip - Type1: Grass, Type2: Flying
4. Spinda - Type1: Normal, Type2:
5. Torterra - Type1: Grass, Type2: Ground
6. Polteageist - Type1: Ghost, Type2:
7. Mismagius - Type1: Ghost, Type2:
8. Teddiursa - Type1: Normal, Type2:
9. Chespin - Type1: Grass, Type2:
10. Dracozolt - Type1: Electric, Type2: Dragon

results_q21 = {}

for pokemon_idx in random_indices:
    pokemon = pokedex.iloc[pokemon_idx]
    pokemon_name = pokemon['Name']

    img_emb = image_embeddings[pokemon_idx:pokemon_idx+1]
    similarities = (img_emb @ text_embeddings_all_types.T).flatten()

    top5_indices = np.argsort(similarities)[::-1][:5]
    top5_types = [all_types[i] for i in top5_indices]
    top5_scores = similarities[top5_indices]

```

```
results_q21[pokemon_idx] = {
    'pokemon': pokemon,
    'top5_types': top5_types,
    'top5_scores': top5_scores
}

print(f"\{pokemon_name} (Type1: {pokemon['Type1']}, Type2: {pokemon['Type2']}):")
for i, (type_name, score) in enumerate(zip(top5_types, top5_scores)):
    print(f" {i+1}. {type_name}: {score:.5f}")

Sawsbuck (Type1: Normal, Type2: Grass):
1. Grass: 0.22229
2. Electric: 0.18726
3. Dark: 0.18542
4. Water: 0.17957
5. Fire: 0.17761
Croagunk (Type1: Poison, Type2: Fighting):
1. Water: 0.21838
2. Fighting: 0.21765
3. Dark: 0.21497
4. Electric: 0.21448
5. Ground: 0.21155
Hoppip (Type1: Grass, Type2: Flying):
1. Grass: 0.23389
2. Poison: 0.22742
3. Fairy: 0.21899
4. Electric: 0.21545
5. Water: 0.20227
Spinda (Type1: Normal, Type2: ):
1. Psychic: 0.20618
2. Electric: 0.19482
3. Ghost: 0.19116
4. Fighting: 0.18665
5. Poison: 0.18555
Torterra (Type1: Grass, Type2: Ground):
1. Rock: 0.22327
2. Grass: 0.22131
3. Water: 0.21130
4. Electric: 0.20679
5. Steel: 0.20654
Polteageist (Type1: Ghost, Type2: ):
1. Water: 0.22095
2. Poison: 0.21594
3. Ghost: 0.20947
4. Psychic: 0.20935
5. Ice: 0.20752
Mismagius (Type1: Ghost, Type2: ):
1. Poison: 0.22473
```

```

2. Ghost: 0.22449
3. Psychic: 0.22095
4. Dark: 0.20850
5. Grass: 0.20410
Teddiursa (Type1: Normal, Type2: ):
1. Grass: 0.20398
2. Fighting: 0.20215
3. Dark: 0.19836
4. Ice: 0.19446
5. Psychic: 0.19263
Chespin (Type1: Grass, Type2: ):
1. Grass: 0.23108
2. Fighting: 0.20288
3. Electric: 0.20117
4. Psychic: 0.19556
5. Ground: 0.19519
Dracozolt (Type1: Electric, Type2: Dragon):
1. Dragon: 0.22058
2. Poison: 0.20300
3. Electric: 0.19995
4. Grass: 0.19824
5. Fighting: 0.19348

# Pics for 10 randomly selected Pokemon with predicted types
fig, axes = plt.subplots(2, 5, figsize=(20, 8))

print("\n--- Pokemon Details ---")

for idx, pokemon_idx in enumerate(random_indices):
    row = idx // 5
    col = idx % 5
    ax = axes[row, col]

    pokemon = pokedex.iloc[pokemon_idx]
    result = results_q21[pokemon_idx]

    img = Image.open(pokemon['image_path'])
    ax.imshow(img)
    ax.axis('off')

    # Only putting the Pokemon name in the image so that image is readable
    # Printing all required details to terminal
    ax.set_title(pokemon['Name'], fontsize=15)
    print(f"\n{idx+1}. {pokemon['Name']} ")
    print(f"  Ground Truth Type1: {pokemon['Type1']} ")
    if pd.notna(pokemon['Type2']) and pokemon['Type2'] != '':
        print(f"  Ground Truth Type2: {pokemon['Type2']} ")
    print("  Top 5 Predicted Types:")
    for i, (type_name, score) in enumerate(zip(result['top5_types'],

```

```
result['top5_scores'])):  
    print(f"    {i+1}. {type_name} ({score:.3f})")  
  
plt.tight_layout()  
plt.savefig('question21.png', dpi=150, bbox_inches='tight')  
plt.show()
```

--- Pokemon Details ---

1. Sawsbuck

Ground Truth Type1: Normal
Ground Truth Type2: Grass
Top 5 Predicted Types:
1. Grass (0.222)
2. Electric (0.187)
3. Dark (0.185)
4. Water (0.180)
5. Fire (0.178)

2. Croagunk

Ground Truth Type1: Poison
Ground Truth Type2: Fighting
Top 5 Predicted Types:
1. Water (0.218)
2. Fighting (0.218)
3. Dark (0.215)
4. Electric (0.214)
5. Ground (0.212)

3. Hoppip

Ground Truth Type1: Grass
Ground Truth Type2: Flying
Top 5 Predicted Types:
1. Grass (0.234)
2. Poison (0.227)
3. Fairy (0.219)
4. Electric (0.215)
5. Water (0.202)

4. Spinda

Ground Truth Type1: Normal
Top 5 Predicted Types:
1. Psychic (0.206)
2. Electric (0.195)
3. Ghost (0.191)
4. Fighting (0.187)
5. Poison (0.186)

5. Torterra

Ground Truth Type1: Grass
Ground Truth Type2: Ground

Top 5 Predicted Types:

1. Rock (0.223)
2. Grass (0.221)
3. Water (0.211)
4. Electric (0.207)
5. Steel (0.207)

6. Polteageist

Ground Truth Type1: Ghost

Top 5 Predicted Types:

1. Water (0.221)
2. Poison (0.216)
3. Ghost (0.209)
4. Psychic (0.209)
5. Ice (0.208)

7. Mismagius

Ground Truth Type1: Ghost

Top 5 Predicted Types:

1. Poison (0.225)
2. Ghost (0.224)
3. Psychic (0.221)
4. Dark (0.208)
5. Grass (0.204)

8. Teddiursa

Ground Truth Type1: Normal

Top 5 Predicted Types:

1. Grass (0.204)
2. Fighting (0.202)
3. Dark (0.198)
4. Ice (0.194)
5. Psychic (0.193)

9. Chespin

Ground Truth Type1: Grass

Top 5 Predicted Types:

1. Grass (0.231)
2. Fighting (0.203)
3. Electric (0.201)
4. Psychic (0.196)
5. Ground (0.195)

10. Dracozolt

Ground Truth Type1: Electric

Ground Truth Type2: Dragon

Top 5 Predicted Types:

1. Dragon (0.221)

2. Poison (0.203)
3. Electric (0.200)
4. Grass (0.198)
5. Fighting (0.193)



Question 22

```

print("Evaluating CLIP on all Pokemon images")
print(f"Total Pokemon: {len(pokedex)}")
print(f"Total types: {len(all_types)}")

all_top5_predictions = []
ground_truth_type1 = []

for pokemon_idx in range(len(pokedex)):
    pokemon = pokedex.iloc[pokemon_idx]

    img_emb = image_embeddings[pokemon_idx:pokemon_idx+1]
    similarities = (img_emb @ text_embeddings_all_types.T).flatten()

    top5_indices = np.argsort(similarities)[::-1][:5]
    top5_types = [all_types[i] for i in top5_indices]

    all_top5_predictions.append(top5_types)
    ground_truth_type1.append(pokemon['Type1'])

acc_at_1 = 0
hit_at_5 = 0

for i in range(len(pokedex)):
    predicted_type1 = all_top5_predictions[i][0]
    true_type1 = ground_truth_type1[i]

    if predicted_type1 == true_type1:

```

```

    acc_at_1 += 1

    if true_type1 in all_top5_predictions[i]:
        hit_at_5 += 1

acc_at_1 /= len(pokedex)
hit_at_5 /= len(pokedex)

print(f"Results:")
print(f"Accuracy@1 (Acc@1): {acc_at_1:.4f} ({acc_at_1*100:.2f}%)")
print(f"Hit@5 (Recall@5): {hit_at_5:.4f} ({hit_at_5*100:.2f}%)")
print(f"Acc@1 counts: {sum(1 for i in range(len(pokedex)) if
all_top5_predictions[i][0] == ground_truth_type1[i])}/{len(pokedex)}")
print(f"Hit@5 counts: {sum(1 for i in range(len(pokedex)) if
ground_truth_type1[i] in all_top5_predictions[i])}/{len(pokedex)}")

Evaluating CLIP on all Pokemon images
Total Pokemon: 751
Total types: 18
Results:
Accuracy@1 (Acc@1): 0.3755 (37.55%)
Hit@5 (Recall@5): 0.7071 (70.71%)
Acc@1 counts: 282/751
Hit@5 counts: 531/751

# Question 23

model_vlm, processor_vlm = load_qwen3_vl()

vlm_reranked_predictions = []

print("Reranking CLIP predictions with VLM")
for pokemon_idx in tqdm(range(len(pokedex))):
    pokemon = pokedex.iloc[pokemon_idx]
    image_path = pokemon['image_path']

    clip_top5 = all_top5_predictions[pokemon_idx]

    candidate_list = ", ".join(clip_top5)
    prompt = f"""Given the Pokemon image, select one most likely
primary type from this candidate list: [{candidate_list}].

Respond with only the type name from the candidate list, nothing
else."""

    vlm_output = qwen_vl_infer_one(model_vlm, processor_vlm,
image_path, prompt, max_new_tokens=32)

    vlm_output_clean = vlm_output.strip().replace(".",",
").replace(",\"", "")
```

```

predicted_type = None
for candidate in clip_top5:
    if candidate.lower() in vlm_output_clean.lower() or
vlm_output_clean.lower() in candidate.lower():
        predicted_type = candidate
        break

if predicted_type is None:
    predicted_type = clip_top5[0]

vlm_reranked_predictions.append(predicted_type)

print(f"VLM processing complete.")

{"model_id": "f92ad0530c8a4042aff66c1955a57a7c", "version_major": 2, "vers
ion_minor": 0}

Reranking CLIP predictions with VLM
100%|██████████| 751/751 [01:49<00:00,  6.89it/s]
VLM processing complete.

# Calculate VLM-reranked Acc@1
vlm_acc_at_1 = 0
for i in range(len(pokedex)):
    if vlm_reranked_predictions[i] == ground_truth_type1[i]:
        vlm_acc_at_1 += 1

vlm_acc_at_1 /= len(pokedex)

print("=*60)
print("Comparison Table")
print("=*60)
print(f"{'Metric':<30} {'Value':<15} {'Percentage':<15}")
print("-*60)
print(f"{'CLIP Acc@1':<30} {acc_at_1:<15.4f} {acc_at_1*100:<15.2f}%")
print(f"{'CLIP Hit@5':<30} {hit_at_5:<15.4f} {hit_at_5*100:<15.2f}%")
print(f"{'VLM-reranked Acc@1':<30} {vlm_acc_at_1:<15.4f}
{vlm_acc_at_1*100:<15.2f}%")
print("=*60)

# Calculate improvement
improvement = vlm_acc_at_1 - acc_at_1
improvement_pct = (improvement / acc_at_1 * 100) if acc_at_1 > 0 else
0
print(f"Improvement: {improvement:+.4f} ({improvement_pct:+.2f}%
relative improvement)")

```

```

# Examples
print("*"*60)
print("Example predictions for the first 10 Pokemon")
print("*"*60)
for i in range(min(10, len(pokedex))):
    pokemon = pokedex.iloc[i]
    print(f"{i+1}. {pokemon['Name']}")
    print(f"  Ground Truth Type1: {ground_truth_type1[i]}")
    if pd.notna(pokemon['Type2']):
        print(f"  Ground Truth Type2: {pokemon['Type2']}")

    print(f"  CLIP Top-1: {all_top5_predictions[i][0]}")
    print(f"  CLIP Top-5: {all_top5_predictions[i]}")
    print(f"  VLM Reranked: {vlm_reranked_predictions[i]}")
    print(f"  CLIP Correct: {'Correct' if all_top5_predictions[i][0] == ground_truth_type1[i] else 'Incorrect'}")
    print(f"  VLM Correct: {'Correct' if vlm_reranked_predictions[i] == ground_truth_type1[i] else 'Incorrect'}")

=====

```

Comparison Table

Metric	Value	Percentage	%
CLIP Acc@1	0.3755	37.55	%
CLIP Hit@5	0.7071	70.71	%
VLM-reranked Acc@1	0.4101	41.01	%

Improvement: +0.0346 (+9.22% relative improvement)

Example predictions for the first 10 Pokemon

1. Bulbasaur

```

Ground Truth Type1: Grass
Ground Truth Type2: Poison
CLIP Top-1: Grass
CLIP Top-5: ['Grass', 'Water', 'Poison', 'Electric', 'Ice']
VLM Reranked: Grass
CLIP Correct: Correct
VLM Correct: Correct

```

2. Ivysaur

```

Ground Truth Type1: Grass
Ground Truth Type2: Poison
CLIP Top-1: Grass
CLIP Top-5: ['Grass', 'Water', 'Poison', 'Fighting', 'Ice']
VLM Reranked: Grass
CLIP Correct: Correct
VLM Correct: Correct

```

3. Charmander

```

Ground Truth Type1: Fire
Ground Truth Type2:

```

```
CLIP Top-1: Fire
CLIP Top-5: ['Fire', 'Water', 'Electric', 'Fighting', 'Grass']
VLM Reranked: Fire
CLIP Correct: Correct
VLM Correct: Correct
4. Charmeleon
Ground Truth Type1: Fire
Ground Truth Type2:
CLIP Top-1: Fire
CLIP Top-5: ['Fire', 'Fighting', 'Water', 'Dragon', 'Electric']
VLM Reranked: Fire
CLIP Correct: Correct
VLM Correct: Correct
5. Squirtle
Ground Truth Type1: Water
Ground Truth Type2:
CLIP Top-1: Water
CLIP Top-5: ['Water', 'Ice', 'Electric', 'Fighting', 'Rock']
VLM Reranked: Water
CLIP Correct: Correct
VLM Correct: Correct
6. Wartortle
Ground Truth Type1: Water
Ground Truth Type2:
CLIP Top-1: Water
CLIP Top-5: ['Water', 'Ice', 'Fighting', 'Electric', 'Rock']
VLM Reranked: Water
CLIP Correct: Correct
VLM Correct: Correct
7. Caterpie
Ground Truth Type1: Bug
Ground Truth Type2:
CLIP Top-1: Poison
CLIP Top-5: ['Poison', 'Grass', 'Water', 'Bug', 'Fighting']
VLM Reranked: Bug
CLIP Correct: Incorrect
VLM Correct: Correct
8. Metapod
Ground Truth Type1: Bug
Ground Truth Type2:
CLIP Top-1: Grass
CLIP Top-5: ['Grass', 'Rock', 'Poison', 'Water', 'Ice']
VLM Reranked: Grass
CLIP Correct: Incorrect
VLM Correct: Incorrect
9. Butterfree
Ground Truth Type1: Bug
Ground Truth Type2: Flying
CLIP Top-1: Fairy
```

```
CLIP Top-5: ['Fairy', 'Bug', 'Flying', 'Poison', 'Normal']
VLM Reranked: Bug
CLIP Correct: Incorrect
VLM Correct: Correct
10. Weedle
Ground Truth Type1: Bug
Ground Truth Type2: Poison
CLIP Top-1: Grass
CLIP Top-5: ['Grass', 'Poison', 'Ground', 'Bug', 'Electric']
VLM Reranked: Ground
CLIP Correct: Incorrect
VLM Correct: Incorrect
```

Did VLM reranking help?

Yes, as shown in the table above, VLM reranking directly improved performance, achieving a 9.22% relative increase in classification accuracy, raising [Acc@1](#) from 37.55% to 41.01%. We see that while CLIP is an effective retriever, as shown by its high 70.71% [Hit@5](#), CLIP lacks the fine-grained reasoning to consistently rank the correct type as its top choice. By using the VLM model to reason over image and text tokens, our system was able to evaluate CLIP's top five candidates and sometimes save the correct label as we see with the Caterpie example above.